

DSC 190


Machine Learning: Representations

Lecture 12 | Part 1

Neural Networks

Recall: Linear Predictor


- **Input:** features $\vec{x} = (x_1, \dots, x_d)^T$
- **Parameters:**
 $\vec{w} = (w_0, w_1, \dots, w_d)^T$
- **Output:** $w_0 + w_1 x_1 + \dots + w_d x_d$



Linear Predictors


- ▶ **Pro:** simple, usually easy to optimize \vec{w}
 - ▶ With square loss, solution given by normal equations
- ▶ **Con:** Decision boundary is linear

Example



Recall: Basis Functions

- **Input:** features \vec{x} , basis functions $\varphi_1, \dots, \varphi_d : \mathbb{R}^d \rightarrow \mathbb{R}$
- **Parameters:**
 $\vec{w} = (w_0, w_1, \dots, w_d)^T$
- **Output:**
 $w_0 + w_1 \varphi_1(\vec{x}) + \dots + w_d \varphi_d(\vec{x})$



Basis Functions

- ▶ **Note:** the basis functions and the weights \vec{w} are **not** chosen at the same time
- ▶ Two step process
- ▶ First, basis functions are chosen and fixed
 - ▶ By hand, by k -means clustering, etc.
- ▶ *Then* the weights \vec{w} are learned

Exercise

Why do this in two steps as opposed to one?

Answer

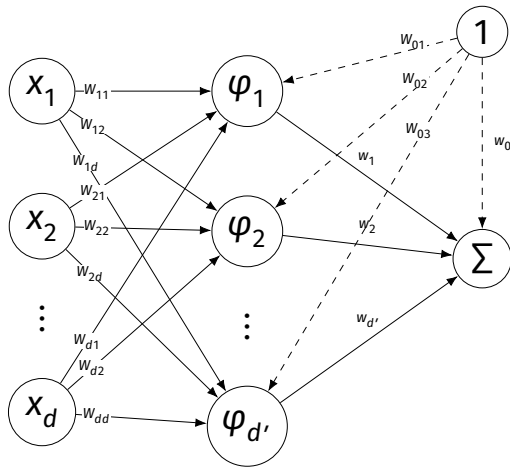
- ▶ By fixing basis functions *then* finding best \vec{w} , optimization is easy again
- ▶ Using square loss, normal equations still work

Idea

- ▶ Try to learn basis functions at same time as weights, \vec{w}
- ▶ Attempt #1: linear basis functions?

$$\varphi_i(\vec{x}) = W_{1i}x_1 + \dots + W_{di}x_d$$


The Model



$$\varphi_i(\vec{x}) = W_{1i}x_1 + \dots + W_{di}x_d$$


Neural Network

- ▶ **Input:** features \vec{x} ,
- ▶ **Parameters:**
 $\vec{w} = (w_0, w_1, \dots, w_d)^T$,
 $(d + 1) \times d'$ matrix W
- ▶ **Output:**
 $w_0 + w_1 \varphi_1(\vec{x}) + \dots + w_d \varphi_d(\vec{x})$
- ▶ This is a **neural network**



Problem


- If φ_i is linear, so is the decision boundary!



Activation Function

- ▶ To make φ_i nonlinear, we often apply a **activation function**.
- ▶ Very commonly: **rectified linear unit** (ReLU)

$$g(z) = \max\{0, z\}$$



$$\begin{aligned}\varphi_i(\vec{x}) &= g(W_{0i} + W_{1i}x_1 + W_{2i}x_2 + \dots + W_{di}x_dA) \\ &= \max\{0, W_{0i} + W_{1i}x_1 + W_{2i}x_2 + \dots + W_{di}x_dA\}\end{aligned}$$


Neural Networks as Functions

- ▶ A neural network is simply a special kind of **function**.
- ▶ $f(\vec{x}; \vec{w}, W)$

Example

$$W = \begin{pmatrix} 2 & -1 \\ 3 & -2 \\ -2 & 1 \end{pmatrix} \quad \vec{w} = \begin{pmatrix} 4 \\ 0 \\ 2 \end{pmatrix} \quad \vec{x} = \begin{pmatrix} 1 \\ 1 \end{pmatrix}$$


The Xor Problem



A Solution

$$W = \begin{pmatrix} 0 & -1 \\ 1 & 1 \\ 1 & 1 \end{pmatrix} \quad \vec{w} = \begin{pmatrix} 0 \\ 1 \\ -2 \end{pmatrix}$$

Prediction Surface



Learning with NNs

- ▶ We can **learn** weights by gathering data, picking a loss function and minimizing loss.
- ▶ The square loss works:

$$R(\vec{w}, W) = \frac{1}{n} \sum_{i=1}^n (f(\vec{x}^{(i)}; \vec{w}, W) - y_i)^2$$

Problem

- ▶ Now that the basis function weights are learnable, too, there is no simple solution for the best weights.
- ▶ We must instead use **gradient descent**.

DSC 190

Machine Learning: Representations

Lecture 12 | Part 2


Gradient Descent

Gradient Descent

- ▶ We have a function $f : \mathbb{R} \rightarrow \mathbb{R}$
- ▶ We can't solve for the x that minimizes (or maximizes) $f(x)$
- ▶ Instead, we use the derivative to “walk” towards the optimizer


Meaning of the Derivative

- ▶ We have the derivative; can we use it?
- ▶ $\frac{df}{dx}(x)$ is a function; it gives the **slope** at x .




Key Idea Behind Gradient Descent

- ▶ If the slope of f at x is **positive** then moving to the **left** decreases the value of f .
- ▶ i.e., we should **decrease** x



Key Idea Behind Gradient Descent

- ▶ If the slope of f at x is **negative** then moving to the **right** decreases the value of f .
- ▶ i.e., we should **increase** x




Key Idea Behind Gradient Descent

- ▶ Pick a starting place, x_0 . Where do we go next?
- ▶ Slope at x_0 negative? Then increase x_0 .
- ▶ Slope at x_0 positive? Then decrease x_0 .
- ▶ This will work:

$$x_1 = x_0 - \frac{df}{dx}(x_0)$$

Gradient Descent

- ▶ Pick α to be a positive number. It is the **learning rate**.
- ▶ Pick a starting prediction, x_0 .
- ▶ On step i , perform update $x_i = x_{i-1} - \alpha \cdot \frac{df}{dx}(x_{i-1})$
- ▶ Repeat until convergence (when x doesn't change much).



```
def gradient_descent(derivative, x, alpha, tol=1e-12):  
    """Minimize using gradient descent."""  
    while True:  
        x_next = x - alpha * derivative(x)  
        if abs(x_next - x) < tol:  
            break  
        x = x_next  
    return h
```

Example: Minimizing Mean Squared Error

- Recall the mean squared error and its derivative:

$$R_{\text{sq}}(x) = \frac{1}{n} \sum_{i=1}^n (x - y_i)^2 \quad \frac{dR_{\text{sq}}}{dx}(x) = \frac{2}{n} \sum_{i=1}^n (x - y_i)$$

Exercise

Let $y_1 = -4$, $y_2 = -2$, $y_3 = 2$, $y_4 = 4$.

Pick $x_0 = 4$ and $\alpha = 1/4$. What is x_1 ?

- a) -1
- b) 0
- c) 1
- d) 2

Example

Gradient Descent in > 1 dimensions

- ▶ The derivative of f becomes the gradient:

$$\frac{df}{dx} \rightarrow \nabla f(\vec{x})$$

- ▶ Meaning of **differentiable**: locally, f looks linear.
- ▶ **Key**: $\nabla f(\vec{w})$ is a function; it returns a vector pointing in direction of steepest ascent.

Gradient Descent in > 1 dimensions

- ▶ Pick α to be a positive number.
 - ▶ It is the **learning rate**.
- ▶ Pick a starting guess, $\vec{w}^{(0)}$.
- ▶ On step i , update $\vec{w}^{(i)} = \vec{w}^{(i-1)} - \alpha \cdot \nabla f(\vec{w}^{(i-1)})$
- ▶ Repeat until convergence
 - ▶ when \vec{w} doesn't change much
 - ▶ equivalently, when $\|\nabla f(\vec{w}^{(i)})\|$ is small


```
def gradient_descent(gradient, w, alpha, tol=1e-12):  
    """Minimize using gradient descent."""  
    while True:  
        w_next = w - alpha * gradient(x)  
        if np.linalg.norm(w_next - w) < tol:  
            break  
        w = w_next  
    return w
```

DSC 190

Machine Learning: Representations


Lecture 12 | Part 3

Convexity in 1-d



Question

When is gradient descent guaranteed to work?

Not here...



Convex Functions




Convexity: Definition

- f is **convex** if for **every** a, b the line segment between

$$(a, f(a)) \quad \text{and} \quad (b, f(b))$$

does not go below the plot of f .




Convexity: Definition

- f is **convex** if for **every** a, b the line segment between

$$(a, f(a)) \quad \text{and} \quad (b, f(b))$$

does not go below the plot of f .




Convexity: Definition

- f is **convex** if for **every** a, b the line segment between

$$(a, f(a)) \quad \text{and} \quad (b, f(b))$$

does not go below the plot of f .




Convexity: Definition

- f is **convex** if for **every** a, b the line segment between

$$(a, f(a)) \quad \text{and} \quad (b, f(b))$$

does not go below the plot of f .




Other Terms

- ▶ If a function is not convex, it is **non-convex**.
- ▶ **Strictly convex**: the line lies strictly above curve.
- ▶ **Concave**: the line lies on or below curve.

Convexity: Formal Definition

- A function $f : \mathbb{R} \rightarrow \mathbb{R}$ is **convex** if for every choice of $a, b \in \mathbb{R}$ and $t \in [0, 1]$:

$$(1 - t)f(a) + tf(b) \geq f((1 - t)a + tb).$$





Example

Is $f(x) = |x|$ convex?

Another View: Second Derivatives

- ▶ If $\frac{d^2f}{dx^2}(x) \geq 0$ for all x , then f is convex.
- ▶ Example: $f(x) = x^4$ is convex.
- ▶ **Warning!** Only works if f is twice differentiable!





Another View: Second Derivatives

- ▶ “Best” straight line at x_0 :
 - ▶ $h_1(z) = f'(x_0) \cdot z + b$
- ▶ “Best” parabola at x_0 :
 - ▶ At x_0 , f looks like $h_2(z) = \frac{1}{2}f''(x_0) \cdot z^2 + f'(x_0)z + c$
 - ▶ Possibilities: upward-facing, downward-facing.

Convexity and Parabolas

- ▶ Convex if for **every** x_0 , parabola is upward-facing.
 - ▶ That is, $f''(x_0) \geq 0$.



Convexity and Gradient Descent

- ▶ Convex functions are (relatively) easy to optimize.
- ▶ **Theorem:** if $R(x)$ is convex and differentiable¹² then gradient descent converges to a **global optimum** of R *provided* that the step size is small enough³.

¹and its derivative is not too wild

²actually, a modified GD works on non-differentiable functions

³step size related to steepness.

Nonconvexity and Gradient Descent

- ▶ Nonconvex functions are (relatively) hard to optimize.
- ▶ Gradient descent can still be useful.
- ▶ But not guaranteed to converge to a global minimum.

DSC 190

Machine Learning: Representations

Lecture 12 | Part 4


Convexity in Many Dimensions

Convexity: Definition

- $f(\vec{x})$ is **convex** if for **every** \vec{a}, \vec{b} the line segment between

$$(\vec{a}, f(\vec{a})) \quad \text{and} \quad (\vec{b}, f(\vec{b}))$$

does not go below the plot of f .



Convexity: Formal Definition

- ▶ A function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is **convex** if for every choice of $\vec{a}, \vec{b} \in \mathbb{R}^d$ and $t \in [0, 1]$:

$$(1 - t)f(\vec{a}) + tf(\vec{b}) \geq f((1 - t)\vec{a} + t\vec{b}).$$

The Second Derivative Test

- ▶ For 1-d functions, convex if second derivative ≥ 0 .
- ▶ For 2-d functions, convex if ???

The Hessian Matrix

- Create the **Hessian** matrix of second derivatives:

$$H(\vec{X}) = \begin{pmatrix} \frac{\partial^2 f}{\partial x_1^2}(\vec{X}) & \frac{\partial^2 f}{\partial x_1 \partial x_2}(\vec{X}) \\ \frac{\partial^2 f}{\partial x_2 \partial x_1}(\vec{X}) & \frac{\partial^2 f}{\partial x_2^2}(\vec{X}) \end{pmatrix}$$

In General

- If $f : \mathbb{R}^d \rightarrow \mathbb{R}$, the **Hessian** at \vec{x} is:

$$H(\vec{x}) = \begin{pmatrix} \frac{\partial^2 f}{\partial x_1^2}(\vec{x}) & \frac{\partial^2 f}{\partial x_1 \partial x_2}(\vec{x}) & \dots & \frac{\partial^2 f}{\partial x_1 \partial x_d}(\vec{x}) \\ \frac{\partial^2 f}{\partial x_2 \partial x_1}(\vec{x}) & \frac{\partial^2 f}{\partial x_2^2}(\vec{x}) & \dots & \frac{\partial^2 f}{\partial x_2 \partial x_d}(\vec{x}) \\ \dots & \dots & \dots & \dots \\ \frac{\partial^2 f}{\partial x_d \partial x_1}(\vec{x}) & \frac{\partial^2 f}{\partial x_d \partial x_2}(\vec{x}) & \dots & \frac{\partial^2 f}{\partial x_d^2}(\vec{x}) \end{pmatrix}$$

The Second Derivative Test

- ▶ A function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is **convex** if for any $\vec{x} \in \mathbb{R}^d$, the Hessian matrix $H(\vec{x})$ is **positive semi-definite**.
- ▶ That is, all eigenvalues are ≥ 0

Next Time

- ▶ Backpropagation and gradient descent for training neural networks.