

DSC 190


Machine Learning: Representations

Lecture 13 | Part 1


Convexity in 1-d

Neural Networks

A NN is just a function: $f(\vec{x}; \vec{w})$



Example



Learning

- ▶ **Given:** a data set $(\vec{x}^{(i)}, y_i)$
- ▶ **Find:** weights \vec{w} minimizing some cost function (e.g., expected square loss):

$$C(\vec{w}) = \frac{1}{n} \sum_{i=1}^n \left(f(\vec{x}^{(i)}; \vec{w}) - y_i \right)^2$$

- ▶ **Problem:** there is no closed-form solution

Gradient Descent


- **Idea:** start at arbitrary $\vec{w}^{(0)}$, walk in direction of gradient:

$$\nabla C = \begin{pmatrix} \frac{\partial C}{\partial w_0} \\ \frac{\partial C}{\partial w_1} \\ \vdots \\ \frac{\partial C}{\partial w_k} \end{pmatrix}$$



Question

When is gradient descent guaranteed to work?

Not here...



Convex Functions




Convexity: Definition

- f is **convex** if for **every** a, b the line segment between

$$(a, f(a)) \quad \text{and} \quad (b, f(b))$$

does not go below the plot of f .




Convexity: Definition

- f is **convex** if for **every** a, b the line segment between

$$(a, f(a)) \quad \text{and} \quad (b, f(b))$$

does not go below the plot of f .




Convexity: Definition

- f is **convex** if for **every** a, b the line segment between

$$(a, f(a)) \quad \text{and} \quad (b, f(b))$$

does not go below the plot of f .




Convexity: Definition

- f is **convex** if for **every** a, b the line segment between

$$(a, f(a)) \quad \text{and} \quad (b, f(b))$$

does not go below the plot of f .




Other Terms

- ▶ If a function is not convex, it is **non-convex**.
- ▶ **Strictly convex**: the line lies strictly above curve.
- ▶ **Concave**: the line lies on or below curve.

Convexity: Formal Definition

- A function $f : \mathbb{R} \rightarrow \mathbb{R}$ is **convex** if for every choice of $a, b \in \mathbb{R}$ and $t \in [0, 1]$:

$$(1 - t)f(a) + tf(b) \geq f((1 - t)a + tb).$$





Example

Is $f(x) = |x|$ convex?

Another View: Second Derivatives

- ▶ If $\frac{d^2f}{dx^2}(x) \geq 0$ for all x , then f is convex.
- ▶ Example: $f(x) = x^4$ is convex.
- ▶ **Warning!** Only works if f is twice differentiable!





Another View: Second Derivatives

- ▶ “Best” straight line at x_0 :
 - ▶ $h_1(z) = f'(x_0) \cdot z + b$
- ▶ “Best” parabola at x_0 :
 - ▶ At x_0 , f looks like $h_2(z) = \frac{1}{2}f''(x_0) \cdot z^2 + f'(x_0)z + c$
 - ▶ Possibilities: upward-facing, downward-facing.

Convexity and Parabolas

- ▶ Convex if for **every** x_0 , parabola is upward-facing.
 - ▶ That is, $f''(x_0) \geq 0$.



Convexity and Gradient Descent

- ▶ Convex functions are (relatively) easy to optimize.
- ▶ **Theorem:** if $R(x)$ is convex and differentiable¹² then gradient descent converges to a **global optimum** of R *provided* that the step size is small enough³.

¹and its derivative is not too wild

²actually, a modified GD works on non-differentiable functions

³step size related to steepness.

Nonconvexity and Gradient Descent

- ▶ Nonconvex functions are (relatively) hard to optimize.
- ▶ Gradient descent can still be useful.
- ▶ But not guaranteed to converge to a global minimum.

DSC 190

Machine Learning: Representations

Lecture 13 | Part 2


Convexity in Many Dimensions

Convexity: Definition

- $f(\vec{x})$ is **convex** if for **every** \vec{a}, \vec{b} the line segment between

$$(\vec{a}, f(\vec{a})) \quad \text{and} \quad (\vec{b}, f(\vec{b}))$$

does not go below the plot of f .



Convexity: Formal Definition

- ▶ A function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is **convex** if for every choice of $\vec{a}, \vec{b} \in \mathbb{R}^d$ and $t \in [0, 1]$:

$$(1 - t)f(\vec{a}) + tf(\vec{b}) \geq f((1 - t)\vec{a} + t\vec{b}).$$

The Second Derivative Test

- ▶ For 1-d functions, convex if second derivative ≥ 0 .
- ▶ For 2-d functions, convex if ???

The Hessian Matrix

- Create the **Hessian** matrix of second derivatives:

$$H(\vec{X}) = \begin{pmatrix} \frac{\partial^2 f}{\partial x_1^2}(\vec{X}) & \frac{\partial^2 f}{\partial x_1 \partial x_2}(\vec{X}) \\ \frac{\partial^2 f}{\partial x_2 \partial x_1}(\vec{X}) & \frac{\partial^2 f}{\partial x_2^2}(\vec{X}) \end{pmatrix}$$

In General

- If $f : \mathbb{R}^d \rightarrow \mathbb{R}$, the **Hessian** at \vec{x} is:

$$H(\vec{x}) = \begin{pmatrix} \frac{\partial^2 f}{\partial x_1^2}(\vec{x}) & \frac{\partial^2 f}{\partial x_1 \partial x_2}(\vec{x}) & \dots & \frac{\partial^2 f}{\partial x_1 \partial x_d}(\vec{x}) \\ \frac{\partial^2 f}{\partial x_2 \partial x_1}(\vec{x}) & \frac{\partial^2 f}{\partial x_2^2}(\vec{x}) & \dots & \frac{\partial^2 f}{\partial x_2 \partial x_d}(\vec{x}) \\ \dots & \dots & \dots & \dots \\ \frac{\partial^2 f}{\partial x_d \partial x_1}(\vec{x}) & \frac{\partial^2 f}{\partial x_d \partial x_2}(\vec{x}) & \dots & \frac{\partial^2 f}{\partial x_d^2}(\vec{x}) \end{pmatrix}$$

The Second Derivative Test

- ▶ A function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is **convex** if for any $\vec{x} \in \mathbb{R}^d$, the Hessian matrix $H(\vec{x})$ is **positive semi-definite**.
- ▶ That is, all eigenvalues are ≥ 0

DSC 190

Machine Learning: Representations

Lecture 13 | Part 3

Basic Backpropagation

Learning

- ▶ **Given:** a data set $(\vec{x}^{(i)}, y_i)$
- ▶ **Find:** weights \vec{w} minimizing some cost function (e.g., expected square loss):

$$C(\vec{w}) = \frac{1}{n} \sum_{i=1}^n (f(\vec{x}^{(i)}; \vec{w}) - y_i)^2$$

- ▶ **Problem:** there is no closed-form solution

Gradient Descent

- **Idea:** start at arbitrary $\vec{w}^{(0)}$, walk in direction of gradient:

$$\nabla C = \begin{pmatrix} \frac{\partial C}{\partial w_0} \\ \frac{\partial C}{\partial w_1} \\ \vdots \\ \frac{\partial C}{\partial w_k} \end{pmatrix}$$

Computing the Gradient

- ▶ To train a neural network, we can use gradient descent.
- ▶ Involves computing the gradient of the cost function.
- ▶ **Backpropagation** is one method for efficiently computing the gradient.

The Gradient

$$\begin{aligned}\nabla_{\vec{w}} C(\vec{w}) &= \nabla_{\vec{w}} \frac{1}{n} \sum_{i=1}^n (f(\vec{x}^{(i)}; \vec{w}) - y_i)^2 \\ &= \frac{1}{n} \sum_{i=1}^n \nabla_{\vec{w}} (f(\vec{x}^{(i)}; \vec{w}) - y_i)^2 \\ &= \frac{1}{n} \sum_{i=1}^n 2 (f(\vec{x}^{(i)}; \vec{w}) - y_i) \nabla_{\vec{w}} f(\vec{x}^{(i)}; \vec{w})\end{aligned}$$

Interpreting the Gradient

$$\nabla_{\vec{w}} C(\vec{w}) = \frac{1}{n} \sum_{i=1}^n 2 (f(\vec{x}^{(i)}; \vec{w}) - y_i) \nabla_{\vec{w}} f(\vec{x}^{(i)}; \vec{w})$$

- ▶ The gradient has one term for each training example, $(\vec{x}^{(i)}, y_i)$
- ▶ If prediction for $\vec{x}^{(i)}$ is good, contribution to gradient is small.
- ▶ $\nabla_{\vec{w}} f(\vec{x}^{(i)}; \vec{w})$ captures how sensitive $f(\vec{x}^{(i)})$ is to value of each parameter.

The Chain Rule

- ▶ Recall the **chain rule** from calculus.

- ▶ Let $f, g : \mathbb{R} \rightarrow \mathbb{R}$

- ▶ Then:


$$\frac{d}{dx}f(g(x)) = f'(g(x)) \cdot g'(x)$$

- ▶ Alternative notation: $\frac{d}{dx}f(g(x)) = \frac{df}{dg} \frac{dg}{dx}(x)$



Example

► $f(x) = x^2; g(x) = 2x + 1$


The Chain Rule for NNs



Computation Graphs



Example




General Formulas

- Derivatives are defined recursively
- Easy to compute derivatives for early layers if we have derivatives for later layers.

$$\frac{\partial f}{\partial w^{(\ell)}} = \frac{\partial f}{\partial a^{(\ell)}} \cdot \frac{\partial a^{(\ell)}}{\partial z^{(\ell)}} \cdot \frac{\partial z^{(\ell)}}{\partial w^{(\ell)}}$$

$$\frac{\partial f}{\partial a^{(\ell)}} = \frac{\partial f}{\partial a^{(\ell+1)}} \cdot \frac{\partial a^{(\ell+1)}}{\partial z^{(\ell+1)}} \cdot \frac{\partial z^{(\ell+1)}}{\partial a^{(\ell)}}$$

- This is **backpropagation**.




Warning

- ▶ The derivatives depend on the network **architecture**
 - ▶ Number of hidden nodes / layers
- ▶ Backprop is done automatically by your NN library

Backpropagation

Compute the derivatives for the last layers first; use them to compute derivatives for earlier layers.



DSC 190


Machine Learning: Representations

Lecture 13 | Part 4



A More Complex Example

Complexity

- The strategy doesn't change much when each layer has more nodes.




Computational Graph



Example

General Formulas



$$\frac{\partial f}{\partial w_{ij}^{(\ell)}} = \frac{\partial f}{\partial a^{(\ell)}} \cdot \frac{\partial a^{(\ell)}}{\partial z^{(\ell)}} \cdot \frac{\partial z^{(\ell)}}{\partial w_{ij}^{(\ell)}}$$

$$\frac{\partial f}{\partial a^{(\ell)}} = \frac{\partial f}{\partial a^{(\ell+1)}} \cdot \frac{\partial a^{(\ell+1)}}{\partial z^{(\ell+1)}} \cdot \frac{\partial z^{(\ell+1)}}{\partial a^{(\ell)}}$$

DSC 190

Machine Learning: Representations

Lecture 13 | Part 5

Intuition Behind Backprop

Intuition

